

Question 1: Minimum Cuts in a Tree

Problem: How many different minimum cuts are there in a tree with n nodes (i.e., $n - 1$ edges)?

Solution: A tree is defined as a connected acyclic graph. A tree with n nodes has exactly $n - 1$ edges.

1. **Min-Cut Size:** By definition, removing any single edge from a tree disconnects the graph into exactly two connected components. Therefore, every edge defines a cut of size 1. Since a connected graph cannot have a cut of size 0, the minimum cut size is exactly 1.
2. **Distinctness:** Each of the $n - 1$ edges connects a unique pair of subtrees. Removing edge e_i creates a partition (A_i, B_i) that is distinct from the partition created by removing any other edge e_j .
3. **Counting:** Since every edge corresponds to a distinct minimum cut, and there are $n - 1$ edges, there are exactly $n - 1$ minimum cuts.

$$\text{Total Min Cuts} = \text{Number of Edges} = n - 1$$

Answer: $n - 1$

Question 2: Probability in Karger's Algorithm

Problem: Let $p = \frac{1}{\binom{n}{2}}$. Analyze the probability $\Pr[\text{out} = (A, B)]$ relative to p .

Solution: The fundamental analysis of Karger's Random Contraction Algorithm provides a lower bound for the success probability. Let k be the size of the minimum cut. Since the min-cut is k , the degree of every vertex is at least k , meaning the graph has at least $nk/2$ edges. The probability that a specific minimum cut (A, B) survives the contraction process is:

$$\Pr[\text{out} = (A, B)] \geq \frac{2}{n(n-1)} = \frac{1}{\binom{n}{2}} = p$$

This inequality holds for **every** minimum cut in **every** graph.

Therefore, the correct statement is: "For every graph G with n nodes and every min cut (A, B) of G , $\Pr[\text{out} = (A, B)] \geq p$."

Question 3: Randomized Selection Probability

Problem: What is the probability that after one iteration of Randomized Select, the subarray size is $\leq \alpha n$ (where $0.5 < \alpha < 1$)?

Solution: Let the sorted elements of the array be z_1, z_2, \dots, z_n . The algorithm picks a pivot p uniformly at random. To ensure the remaining subarray has size at most αn , the pivot must partition the array such that neither the left side nor the right side is larger than αn . This occurs if the pivot falls in the "middle" of the sorted sequence.

- **Bad Pivot (Left):** If p is among the smallest $(1 - \alpha)n$ elements, the right partition will be larger than αn .
- **Bad Pivot (Right):** If p is among the largest $(1 - \alpha)n$ elements, the left partition will be larger than αn .

The Good Pivot range is the indices between $(1 - \alpha)n$ and αn .

$$\text{Size of Good Range} = \alpha n - (1 - \alpha)n = (2\alpha - 1)n = 2\alpha n - n$$

Since the pivot is chosen uniformly from n elements:

$$Pr(\text{Good Pivot}) = \frac{(2\alpha - 1)n}{n} = 2\alpha - 1$$

Answer: $2\alpha - 1$

Question 4: Recursion Depth of RSelect

Problem: If every recursive call reduces the input size to at most an α fraction ($0 < \alpha < 1$), what is the maximum recursion depth?

Solution: Let N_k be the input size at depth k . We are given the recurrence inequality:

$$N_k \leq \alpha \cdot N_{k-1}$$

Unrolling this from the initial size n :

$$N_k \leq n \cdot \alpha^k$$

The recursion terminates when the problem size reduces to 1 (base case).

$$1 = n \cdot \alpha^k$$

Taking the logarithm of both sides:

$$\log(1) = \log(n \cdot \alpha^k)$$

$$0 = \log(n) + k \log(\alpha)$$

$$-k \log(\alpha) = \log(n)$$

$$k = -\frac{\log(n)}{\log(\alpha)}$$

(Note: Since $\alpha < 1$, $\log(\alpha)$ is negative, so the depth is positive).

Answer: $-\frac{\log n}{\log \alpha}$

Question 5: Global Min Cut via s-t Min Cut

Problem: Minimum number of calls to an s-t min cut subroutine needed to find the global min cut.

Solution: The global minimum cut partitions the vertices V into two disjoint sets, S and $V \setminus S$.

1. We fix an arbitrary vertex s . This vertex must belong to one of the two partitions (without loss of generality, assume $s \in S$).
2. The other partition $V \setminus S$ is non-empty, so there exists at least one vertex $t \in V \setminus S$.
3. We iterate through all possible candidate vertices for t (all $v \in V, v \neq s$).
4. By computing the Minimum $s - t$ Cut for every $t \in V \setminus \{s\}$, we are guaranteed that at least one iteration will have t in the correct opposing set ($V \setminus S$).
5. For that specific pair, the min $s - t$ cut will correspond exactly to the global min cut.

Since we fix s and iterate through all other nodes, the number of calls is $n - 1$.

Answer: $n - 1$

Optional Problem 1: Lower Bound for Randomized Sorting

Problem: Prove that the worst-case expected running time of every randomized comparison-based sorting algorithm is $\Omega(n \log n)$.

Solution: To prove this, we utilize **Yao's Minimax Principle**, which allows us to establish a lower bound on the worst-case expected cost of a randomized algorithm by analyzing the average-case cost of a deterministic algorithm over a specific input distribution.

Key Idea: The average performance of a randomized algorithm on its single worst-case input ... is the same as ... The average performance of the best possible deterministic algorithm on a randomly shuffled input.

1. Definitions:

- Let \mathcal{A} be a randomized comparison-based sorting algorithm. We can view \mathcal{A} as a probability distribution over a set of deterministic algorithms $\{D_1, D_2, \dots\}$.
- Let $T(D, I)$ be the running time (number of comparisons) of deterministic algorithm D on input I .
- We are looking for the worst-case expected running time: $\max_I E[\mathcal{T}(\mathcal{A}, I)]$.

2. Yao's Minimax Principle:

Yao's principle states that for any probability distribution \mathcal{P} over the set of inputs:

$$\max_I E_{\mathcal{A}}[T(\mathcal{A}, I)] \geq \min_D E_{I \sim \mathcal{P}}[T(D, I)]$$

In words: The worst-case expected time of the best randomized algorithm is at least the average-case time of the best deterministic algorithm for the "worst" input distribution.

3. Choosing the Input Distribution:

Let \mathcal{P} be the uniform distribution over all $n!$ possible permutations of the input array.

4. Deterministic Lower Bound:

Consider any deterministic comparison-based sorting algorithm D . It can be modeled as a decision tree.

- The tree must have at least $n!$ leaves (one for each permutation) to correctly distinguish all inputs.
- For a binary tree with L leaves, the average depth (average path length from root to leaf) is lower bounded by $\Omega(\log L)$.
- Substituting $L = n!$, the average depth is $\Omega(\log(n!))$.
- Using Stirling's approximation ($\log(n!) = \Theta(n \log n)$), the average running time of any deterministic algorithm D under the uniform distribution is $\Omega(n \log n)$.

5. Conclusion:

Since $\min_D E_{I \sim \mathcal{P}}[T(D, I)] = \Omega(n \log n)$, by Yao's principle, the worst-case expected running time of any randomized algorithm is also $\Omega(n \log n)$.

Optional Problem 2: Deterministic Selection with Different Group Sizes

Problem: Analyze the running time of the deterministic selection algorithm (Median-of-Medians) if we use group sizes of 7 and 3.

Part A: Group Size of 7

Let $T(n)$ be the running time on an input of size n . The algorithm steps are:

1. Divide elements into $\lceil n/7 \rceil$ groups of 7. Find the median of each group. (Time: $O(n)$)
2. Recursively find the median of these medians, call it x . (Time: $T(n/7)$)
3. Partition the original array around x . (Time: $O(n)$)
4. Recurse on the appropriate subarray. (Time: $T(k)$, where k is the max subarray size)

Calculating the Recurrence:

- The number of groups is $\approx \frac{n}{7}$.
- The pivot x is the median of the $\frac{n}{7}$ group medians.
- Therefore, x is greater than $\frac{1}{2} \cdot \frac{n}{7} = \frac{n}{14}$ group medians.
- In each of these $\frac{n}{14}$ groups, there are 4 elements (the median itself and 3 elements greater than it) that are definitely $\geq x$.
- Total elements guaranteed $\geq x$: $4 \times \frac{n}{14} = \frac{2n}{7}$.
- Similarly, total elements guaranteed $\leq x$: $\frac{2n}{7}$.
- In the worst case, the algorithm recurses on the remaining elements: $n - \frac{2n}{7} = \frac{5n}{7}$.

The recurrence is:

$$T(n) \leq T\left(\frac{n}{7}\right) + T\left(\frac{5n}{7}\right) + O(n)$$

We check if the sum of fractions is less than 1:

$$\frac{1}{7} + \frac{5}{7} = \frac{6}{7} < 1$$

Since the sum is strictly less than 1, the running time is **linear**, i.e., $O(n)$.

Part B: Group Size of 3

Using the same logic for groups of 3:

- Number of groups: $\frac{n}{3}$.
- Pivot x is greater than $\frac{1}{2} \cdot \frac{n}{3} = \frac{n}{6}$ medians.
- In each such group, there are 2 elements (median and 1 greater) definitely $\geq x$.
- Total elements guaranteed $\geq x$: $2 \times \frac{n}{6} = \frac{n}{3}$.
- Max elements in recursive step: $n - \frac{n}{3} = \frac{2n}{3}$.

The recurrence is:

$$T(n) \leq T\left(\frac{n}{3}\right) + T\left(\frac{2n}{3}\right) + O(n)$$

Sum of fractions:

$$\frac{1}{3} + \frac{2}{3} = 1$$

Since the sum equals 1, this recurrence behaves like Merge Sort (specifically $T(n) = T(n/3) + T(2n/3) + cn$). The work per level is $O(n)$, and there are $\log n$ levels. Thus, the running time is $O(n \log n)$, not linear.

Optional Problem 3: Linear Time Weighted Median

Problem: Compute all weighted medians in $O(n)$ worst-case time.

Solution:

We can adapt the standard deterministic selection algorithm. The key insight is that we can discard a portion of the array if we know the weighted median cannot be inside it, provided we account for the weights of the discarded elements.

Algorithm: `WeightedSelect(Array A, TargetWeight T)`

1. **Base Case:** If $|A|$ is small (e.g., 1 or 2), solve by brute force.
2. **Pivot Selection:** Use the deterministic Median-of-Medians algorithm to find the *unweighted* median of A , call it p . This takes $O(|A|)$ time.
3. **Partition:** Partition A around p into three sets:
 - $L = \{x \in A \mid x < p\}$
 - $E = \{x \in A \mid x = p\}$
 - $G = \{x \in A \mid x > p\}$
4. **Compute Weights:** Calculate the total weight of each set: $W_L = \sum_{x \in L} w_x$, $W_E = \sum_{x \in E} w_x$, $W_G = \sum_{x \in G} w_x$.
5. **Recurse:**
 - **Case 1:** If $W_L > T$, the weighted median must be in L .
Recurse: `WeightedSelect(L, T)`
 - **Case 2:** If $W_L + W_E \geq T$, then p is the weighted median (or one of them).
Return p .
 - **Case 3:** If $W_L + W_E < T$, the weighted median is in G . However, we effectively "removed" $W_L + W_E$ weight from the left. We must adjust the target.
Recurse: `WeightedSelect(G, T - (W_L + W_E))`

Complexity Analysis: Since we use the median-of-medians as the pivot p , we guarantee that $|L| \leq \frac{7}{10}n$ and $|G| \leq \frac{7}{10}n$ (approx). The recurrence is:

$$T(n) \leq T(0.7n) + O(n)$$

This geometric series sums to $O(n)$. Thus, the algorithm runs in **O(n)** worst-case time. To handle the initial call, we set $T = W_{total}/2$.

Optional Problem 4: Minimum Cuts in Directed Graphs

Problem: Is it true that every directed graph has only polynomially different minimum cuts? Prove it or give a counterexample.

Solution:

No, this is not true for directed graphs. A directed graph can have exponentially many global minimum cuts.

Counterexample: Consider a simple directed graph G with $n = k + 2$ vertices: a source s , a sink t , and k intermediate nodes v_1, v_2, \dots, v_k .

- Add edges $s \rightarrow v_i$ with capacity 1 for all $i = 1 \dots k$.
- Add edges $v_i \rightarrow t$ with capacity 1 for all $i = 1 \dots k$.
- Add a back-edge $t \rightarrow s$ with capacity ∞ .

Analysis:

1. A global minimum cut partitions vertices into (S, \bar{S}) .
2. Due to the infinite edge $t \rightarrow s$, we cannot choose a cut where $t \in S$ and $s \in \bar{S}$.
3. Thus, valid finite cuts must have $s \in S$ and $t \in \bar{S}$.
4. For every intermediate node v_i , we have a choice:
 - Put $v_i \in S$: The edge $v_i \rightarrow t$ is cut (cost 1).
 - Put $v_i \in \bar{S}$: The edge $s \rightarrow v_i$ is cut (cost 1).
5. Regardless of the choice for each v_i , the cut size is exactly $1 \times k = k$.
6. Since there are k intermediate nodes and 2 choices for each, there are 2^k distinct cuts of size k .
7. Since $k = n - 2$, the number of minimum cuts is 2^{n-2} , which is **exponential** in n .

Optional Problem 5: Number of α -Minimum Cuts

Problem: How many α -minimum cuts can an undirected graph have? Prove the best upper bound.

Solution:

The upper bound is $O(n^{2\alpha})$.

Proof using Karger's Algorithm: Let k be the size of the minimum cut. An α -minimum cut has size at most αk .

1. In a graph with min-cut k , the minimum degree is at least k . Thus, the total number of edges $m \geq \frac{nk}{2}$.
2. Consider a specific α -minimum cut C with size $|C| \leq \alpha k$.
3. In Karger's contraction algorithm, when the graph has r vertices remaining, the probability of choosing an edge in C to contract is:

$$P(\text{edge in } C) = \frac{|C|}{\text{Edges remaining}} \leq \frac{\alpha k}{rk/2} = \frac{2\alpha}{r}$$

4. The probability that C survives this step is $1 - \frac{2\alpha}{r}$.

5. The algorithm stops contraction when 2α vertices remain (generalized version). The probability that C survives from n vertices down to 2α vertices is:

$$P(\text{survive}) \geq \prod_{r=2\alpha+1}^n \left(1 - \frac{2\alpha}{r}\right)$$

6. This product approximates to:

$$\approx \frac{\binom{2\alpha}{2\alpha}}{\binom{n}{2\alpha}} = \frac{1}{\binom{n}{2\alpha}} \approx n^{-2\alpha}$$

7. Since each run of the algorithm produces the specific cut C with probability at least $n^{-2\alpha}$, and the events for different cuts are disjoint in a single run output, the maximum number of such cuts is the reciprocal of the probability.

Result: There are at most $O(n^{2\alpha})$ cuts of size $\alpha \times \text{min-cut}$.