

## Question 1: Minimum Cuts in a Tree

**Problem:** How many different minimum cuts are there in a tree with  $n$  nodes (i.e.,  $n - 1$  edges)?

**Solution:** A tree is defined as a connected acyclic graph. A tree with  $n$  nodes has exactly  $n - 1$  edges.

1. **Min-Cut Size:** By definition, removing any single edge from a tree disconnects the graph into exactly two connected components. Therefore, every edge defines a cut of size 1. Since a connected graph cannot have a cut of size 0, the minimum cut size is exactly 1.
2. **Distinctness:** Each of the  $n - 1$  edges connects a unique pair of subtrees. Removing edge  $e_i$  creates a partition  $(A_i, B_i)$  that is distinct from the partition created by removing any other edge  $e_j$ .
3. **Counting:** Since every edge corresponds to a distinct minimum cut, and there are  $n - 1$  edges, there are exactly  $n - 1$  minimum cuts.

$$\text{Total Min Cuts} = \text{Number of Edges} = n - 1$$

**Answer:**  $n - 1$

---

## Question 2: Probability in Karger's Algorithm

**Problem:** Let  $p = \frac{1}{\binom{n}{2}}$ . Analyze the probability  $\Pr[\text{out} = (A, B)]$  relative to  $p$ .

**Solution:** The fundamental analysis of Karger's Random Contraction Algorithm provides a lower bound for the success probability. Let  $k$  be the size of the minimum cut. Since the min-cut is  $k$ , the degree of every vertex is at least  $k$ , meaning the graph has at least  $nk/2$  edges. The probability that a specific minimum cut  $(A, B)$  survives the contraction process is:

$$\Pr[\text{out} = (A, B)] \geq \frac{2}{n(n-1)} = \frac{1}{\binom{n}{2}} = p$$

This inequality holds for **every** minimum cut in **every** graph.

Therefore, the correct statement is: "For every graph  $G$  with  $n$  nodes and every min cut  $(A, B)$  of  $G$ ,  $\Pr[\text{out} = (A, B)] \geq p$ ."

---

## Question 3: Randomized Selection Probability

**Problem:** What is the probability that after one iteration of Randomized Select, the subarray size is  $\leq \alpha n$  (where  $0.5 < \alpha < 1$ )?

**Solution:** Let the sorted elements of the array be  $z_1, z_2, \dots, z_n$ . The algorithm picks a pivot  $p$  uniformly at random. To ensure the remaining subarray has size at most  $\alpha n$ , the pivot must partition the array such that neither the left side nor the right side is larger than  $\alpha n$ . This occurs if the pivot falls in the "middle" of the sorted sequence.

- **Bad Pivot (Left):** If  $p$  is among the smallest  $(1 - \alpha)n$  elements, the right partition will be larger than  $\alpha n$ .
- **Bad Pivot (Right):** If  $p$  is among the largest  $(1 - \alpha)n$  elements, the left partition will be larger than  $\alpha n$ .

The "Good Pivot" range is the indices between  $(1 - \alpha)n$  and  $\alpha n$ .

$$\text{Size of Good Range} = \alpha n - (1 - \alpha)n = (2\alpha - 1)n$$

Since the pivot is chosen uniformly from  $n$  elements:

$$Pr(\text{Good Pivot}) = \frac{(2\alpha - 1)n}{n} = 2\alpha - 1$$

**Answer:**  $2\alpha - 1$

---

## Question 4: Recursion Depth of RSelect

**Problem:** If every recursive call reduces the input size to at most an  $\alpha$  fraction ( $0 < \alpha < 1$ ), what is the maximum recursion depth?

**Solution:** Let  $N_k$  be the input size at depth  $k$ . We are given the recurrence inequality:

$$N_k \leq \alpha \cdot N_{k-1}$$

Unrolling this from the initial size  $n$ :

$$N_k \leq n \cdot \alpha^k$$

The recursion terminates when the problem size reduces to 1 (base case).

$$1 = n \cdot \alpha^k$$

Taking the logarithm of both sides:

$$\log(1) = \log(n \cdot \alpha^k)$$

$$0 = \log(n) + k \log(\alpha)$$

$$-k \log(\alpha) = \log(n)$$

$$k = -\frac{\log(n)}{\log(\alpha)}$$

(Note: Since  $\alpha < 1$ ,  $\log(\alpha)$  is negative, so the depth is positive).

**Answer:**  $-\frac{\log n}{\log \alpha}$

---

## Question 5: Global Min Cut via s-t Min Cut

**Problem:** Minimum number of calls to an s-t min cut subroutine needed to find the global min cut.

**Solution:** The global minimum cut partitions the vertices  $V$  into two disjoint sets,  $S$  and  $V \setminus S$ .

1. We fix an arbitrary vertex  $s$ . This vertex must belong to one of the two partitions (without loss of generality, assume  $s \in S$ ).
2. The other partition  $V \setminus S$  is non-empty, so there exists at least one vertex  $t \in V \setminus S$ .
3. We iterate through all possible candidate vertices for  $t$  (all  $v \in V, v \neq s$ ).
4. By computing the Minimum  $s - t$  Cut for every  $t \in V \setminus \{s\}$ , we are guaranteed that at least one iteration will have  $t$  in the correct opposing set ( $V \setminus S$ ).
5. For that specific pair, the min  $s - t$  cut will correspond exactly to the global min cut.

Since we fix  $s$  and iterate through all other nodes, the number of calls is  $n - 1$ .

**Answer:**  $n - 1$

## Optional Problem 1: Lower Bound for Randomized Sorting

**Problem:** Prove that the worst-case expected running time of every randomized comparison-based sorting algorithm is  $\Omega(n \log n)$ .

**Solution:**

To prove this, we utilize **Yao's Minimax Principle**, which allows us to establish a lower bound on the worst-case expected cost of a randomized algorithm by analyzing the average-case cost of a deterministic algorithm over a specific input distribution.

**1. Definitions:**

- Let  $\mathcal{A}$  be a randomized comparison-based sorting algorithm. We can view  $\mathcal{A}$  as a probability distribution over a set of deterministic algorithms  $\{D_1, D_2, \dots\}$ .
- Let  $T(D, I)$  be the running time (number of comparisons) of deterministic algorithm  $D$  on input  $I$ .
- We are looking for the worst-case expected running time:  $\max_I E[\mathcal{T}(\mathcal{A}, I)]$ .

**2. Yao's Minimax Principle:** Yao's principle states that for any probability distribution  $\mathcal{P}$  over the set of inputs:

$$\max_I E_{\mathcal{A}}[T(\mathcal{A}, I)] \geq \min_D E_{I \sim \mathcal{P}}[T(D, I)]$$

In words: The worst-case expected time of the best randomized algorithm is at least the average-case time of the best deterministic algorithm for the "worst" input distribution.

**3. Choosing the Input Distribution:** Let  $\mathcal{P}$  be the uniform distribution over all  $n!$  possible permutations of the input array.

**4. Deterministic Lower Bound:** Consider any deterministic comparison-based sorting algorithm  $D$ . It can be modeled as a decision tree.

- The tree must have at least  $n!$  leaves (one for each permutation) to correctly distinguish all inputs.
- For a binary tree with  $L$  leaves, the average depth (average path length from root to leaf) is lower bounded by  $\Omega(\log L)$ .
- Substituting  $L = n!$ , the average depth is  $\Omega(\log(n!))$ .
- Using Stirling's approximation ( $\log(n!) = \Theta(n \log n)$ ), the average running time of any deterministic algorithm  $D$  under the uniform distribution is  $\Omega(n \log n)$ .

**5. Conclusion:** Since  $\min_D E_{I \sim \mathcal{P}}[T(D, I)] = \Omega(n \log n)$ , by Yao's principle, the worst-case expected running time of any randomized algorithm is also  $\Omega(n \log n)$ .

## Optional Problem 2: Deterministic Selection with Different Group Sizes

**Problem:** Analyze the running time of the deterministic selection algorithm (Median-of-Medians) if we use group sizes of 7 and 3.

### Part A: Group Size of 7

Let  $T(n)$  be the running time on an input of size  $n$ . The algorithm steps are:

1. Divide elements into  $\lceil n/7 \rceil$  groups of 7. Find the median of each group. (Time:  $O(n)$ )
2. Recursively find the median of these medians, call it  $x$ . (Time:  $T(n/7)$ )

3. Partition the original array around  $x$ . (Time:  $O(n)$ )
4. Recurse on the appropriate subarray. (Time:  $T(k)$ , where  $k$  is the max subarray size)

**Calculating the Recurrence:**

- The number of groups is  $\approx \frac{n}{7}$ .
- The pivot  $x$  is the median of the  $\frac{n}{7}$  group medians.
- Therefore,  $x$  is greater than  $\frac{1}{2} \cdot \frac{n}{7} = \frac{n}{14}$  group medians.
- In each of these  $\frac{n}{14}$  groups, there are 4 elements (the median itself and 3 elements greater than it) that are definitely  $\geq x$ .
- Total elements guaranteed  $\geq x$ :  $4 \times \frac{n}{14} = \frac{2n}{7}$ .
- Similarly, total elements guaranteed  $\leq x$ :  $\frac{2n}{7}$ .
- In the worst case, the algorithm recurses on the remaining elements:  $n - \frac{2n}{7} = \frac{5n}{7}$ .

The recurrence is:

$$T(n) \leq T\left(\frac{n}{7}\right) + T\left(\frac{5n}{7}\right) + O(n)$$

We check if the sum of fractions is less than 1:

$$\frac{1}{7} + \frac{5}{7} = \frac{6}{7} < 1$$

Since the sum is strictly less than 1, the running time is **linear**, i.e.,  $O(n)$ .

## Part B: Group Size of 3

Using the same logic for groups of 3:

- Number of groups:  $\frac{n}{3}$ .
- Pivot  $x$  is greater than  $\frac{1}{2} \cdot \frac{n}{3} = \frac{n}{6}$  medians.
- In each such group, there are 2 elements (median and 1 greater) definitely  $\geq x$ .
- Total elements guaranteed  $\geq x$ :  $2 \times \frac{n}{6} = \frac{n}{3}$ .
- Max elements in recursive step:  $n - \frac{n}{3} = \frac{2n}{3}$ .

The recurrence is:

$$T(n) \leq T\left(\frac{n}{3}\right) + T\left(\frac{2n}{3}\right) + O(n)$$

Sum of fractions:

$$\frac{1}{3} + \frac{2}{3} = 1$$

Since the sum equals 1, this recurrence behaves like Merge Sort (specifically  $T(n) = T(n/3) + T(2n/3) + cn$ ). The work per level is  $O(n)$ , and there are  $\log n$  levels. Thus, the running time is  $O(n \log n)$ , not linear.

## Optional Problem 3: Linear Time Weighted Median

**Problem:** Compute all weighted medians in  $O(n)$  worst-case time.

**Solution:**

We can adapt the standard deterministic selection algorithm. The key insight is that we can discard a portion of the array if we know the weighted median cannot be inside it, provided we account for the weights of the discarded elements.

**Algorithm:** `WeightedSelect(Array A, TargetWeight T)`

1. **Base Case:** If  $|A|$  is small (e.g., 1 or 2), solve by brute force.
2. **Pivot Selection:** Use the deterministic Median-of-Medians algorithm to find the *unweighted* median of  $A$ , call it  $p$ . This takes  $O(|A|)$  time.
3. **Partition:** Partition  $A$  around  $p$  into three sets:
  - $L = \{x \in A \mid x < p\}$
  - $E = \{x \in A \mid x = p\}$
  - $G = \{x \in A \mid x > p\}$
4. **Compute Weights:** Calculate the total weight of each set:  $W_L = \sum_{x \in L} w_x$ ,  $W_E = \sum_{x \in E} w_x$ ,  $W_G = \sum_{x \in G} w_x$ .
5. **Recurse:**
  - **Case 1:** If  $W_L > T$ , the weighted median must be in  $L$ .  
**Recurse:** `WeightedSelect(L, T)`
  - **Case 2:** If  $W_L + W_E \geq T$ , then  $p$  is the weighted median (or one of them).  
**Return**  $p$ .
  - **Case 3:** If  $W_L + W_E < T$ , the weighted median is in  $G$ . However, we effectively "removed"  $W_L + W_E$  weight from the left. We must adjust the target.  
**Recurse:** `WeightedSelect(G, T - (W_L + W_E))`

**Complexity Analysis:** Since we use the median-of-medians as the pivot  $p$ , we guarantee that  $|L| \leq \frac{7}{10}n$  and  $|G| \leq \frac{3}{10}n$  (approx). The recurrence is:

$$T(n) \leq T(0.7n) + O(n)$$

This geometric series sums to  $O(n)$ . Thus, the algorithm runs in **O(n)** worst-case time. To handle the initial call, we set  $T = W_{total}/2$ .

## Optional Problem 4: Minimum Cuts in Directed Graphs

**Problem:** Is it true that every directed graph has only polynomially many different minimum cuts? Prove it or give a counterexample.

**Solution:**

No, this is not true for directed graphs. A directed graph can have exponentially many global minimum cuts.

**Counterexample:** Consider a simple directed graph  $G$  with  $n = k + 2$  vertices: a source  $s$ , a sink  $t$ , and  $k$  intermediate nodes  $v_1, v_2, \dots, v_k$ .

- Add edges  $s \rightarrow v_i$  with capacity 1 for all  $i = 1 \dots k$ .
- Add edges  $v_i \rightarrow t$  with capacity 1 for all  $i = 1 \dots k$ .
- Add a back-edge  $t \rightarrow s$  with capacity  $\infty$ .

### Analysis:

1. A global minimum cut partitions vertices into  $(S, \bar{S})$ .
  2. Due to the infinite edge  $t \rightarrow s$ , we cannot choose a cut where  $t \in S$  and  $s \in \bar{S}$ .
  3. Thus, valid finite cuts must have  $s \in S$  and  $t \in \bar{S}$ .
  4. For every intermediate node  $v_i$ , we have a choice:
    - Put  $v_i \in S$ : The edge  $v_i \rightarrow t$  is cut (cost 1).
    - Put  $v_i \in \bar{S}$ : The edge  $s \rightarrow v_i$  is cut (cost 1).
  5. Regardless of the choice for each  $v_i$ , the cut size is exactly  $1 \times k = k$ .
  6. Since there are  $k$  intermediate nodes and 2 choices for each, there are  $2^k$  distinct cuts of size  $k$ .
  7. Since  $k = n - 2$ , the number of minimum cuts is  $2^{n-2}$ , which is **exponential** in  $n$ .
- 

## Optional Problem 5: Number of $\alpha$ -Minimum Cuts

**Problem:** How many  $\alpha$ -minimum cuts can an undirected graph have? Prove the best upper bound.

**Solution:**

The upper bound is  $O(n^{2\alpha})$ .

**Proof using Karger's Algorithm:** Let  $k$  be the size of the minimum cut. An  $\alpha$ -minimum cut has size at most  $\alpha k$ .

1. In a graph with min-cut  $k$ , the minimum degree is at least  $k$ . Thus, the total number of edges  $m \geq \frac{nk}{2}$ .
2. Consider a specific  $\alpha$ -minimum cut  $C$  with size  $|C| \leq \alpha k$ .
3. In Karger's contraction algorithm, when the graph has  $r$  vertices remaining, the probability of choosing an edge in  $C$  to contract is:

$$P(\text{edge in } C) = \frac{|C|}{\text{Edges remaining}} \leq \frac{\alpha k}{rk/2} = \frac{2\alpha}{r}$$

4. The probability that  $C$  survives this step is  $1 - \frac{2\alpha}{r}$ .
5. The algorithm stops contraction when  $2\alpha$  vertices remain (generalized version). The probability that  $C$  survives from  $n$  vertices down to  $2\alpha$  vertices is:

$$P(\text{survive}) \geq \prod_{r=2\alpha+1}^n \left(1 - \frac{2\alpha}{r}\right)$$

6. This product approximates to:

$$\approx \frac{\binom{2\alpha}{2\alpha}}{\binom{n}{2\alpha}} = \frac{1}{\binom{n}{2\alpha}} \approx n^{-2\alpha}$$

7. Since each run of the algorithm produces the specific cut  $C$  with probability at least  $n^{-2\alpha}$ , and the events for different cuts are disjoint in a single run output, the maximum number of such cuts is the reciprocal of the probability.

**Result:** There are at most  $O(n^{2\alpha})$  cuts of size  $\alpha \times \text{min-cut}$ .