

Exam Review: Problem Set 5 & Red-Black Trees

Part 1: Priority Queues & Heaps

Question 3: Sorted Array Implementation

Scenario: Implementing a Priority Queue using a sorted array (largest to smallest).

- **Extract-Min:** The minimum element is at the end of the array (index $n - 1$). Removing it takes constant time.
- **Insert:** We must use Binary Search to find the position ($O(\log n)$), but then we must shift all smaller elements to make space. In the worst case (inserting a new maximum), this takes linear time.

Answer: Insert is $\Theta(n)$ and Extract-Min is $\Theta(1)$.

Question 4: Unsorted Array Implementation

Scenario: Implementing a Priority Queue using an unsorted array.

- **Insert:** Order does not matter, so we can simply append the new element to the end.
- **Extract-Min:** The minimum could be anywhere. We must scan the entire array (n elements) to find it.

Answer: Insert is $\Theta(1)$ and Extract-Min is $\Theta(n)$.

Question 5: Heap Capabilities

Scenario: Binary Heap with n elements. What can be done in $O(\log n)$?

- **Find largest:** In a Min-Heap, the largest element is at a leaf. We might need to check all $\lceil n/2 \rceil$ leaves. Time: $\Theta(n)$.
- **Find median:** Heaps are not fully sorted. Finding the median generally takes $\Theta(n)$ or $\Theta(n \log n)$.
- **Find fifth-smallest:** The smallest is the root. The 2nd and 3rd are in level 2. The 4th and 5th are near the top. Since $k = 5$ is a constant, searching the top few levels takes constant time relative to n .

Answer: Find the fifth-smallest element.

Part 2: Shortest Paths (Graphs)

Question 1: Path Properties

Scenario: Directed graph with distinct, non-negative edge lengths.

- **True:** The shortest path might have $n - 1$ edges (e.g., a line graph).
- **True:** There is a shortest path with no repeated vertices (a simple path). Cycles with non-negative weights can be removed without increasing path length.
- **False:** It must include the minimum-length edge (it might not be on the path).
- **False:** It must exclude the maximum-length edge (it might be a necessary bridge).

Answer: Statements 1 and 2 are true.

Question 2: Modified Dijkstra

Scenario: Edges leaving source s can be negative. No edges enter s . All other edges non-negative.

- Dijkstra normally fails with negative edges because it assumes a "closed" node's distance is final.
- Here, negative edges only occur at the start. Since no edges return to s , we process neighbors of s first.
- We can view this as adding a constant M to all edges leaving s to make them positive. This increases all path lengths by exactly M , preserving the relative order of shortest paths.

Answer: Dijkstra always works in this specific case.

Part 3: Red-Black Trees

Re-implementation Requirements

To maintain Red-Black invariants (no double reds, equal black height):

- **Search:** No change needed (standard BST search).
- **Insert:** Requires recoloring and rotations to fix violations (e.g., Red parent + Red child).
- **Delete:** Removing a node can violate Black Height. Requires complex fix-up (rotations/recoloring).

Answer: Insert and Delete must be re-implemented.

Height Guarantee

The height of a Red-Black tree is $\Theta(\log n)$ because:

1. Every root-NULL path has $\leq \log_2(n + 1)$ black nodes.
2. No red node has a red child, so at most 50% of nodes on a path are red.
3. Therefore, total height $\leq 2 \times (\text{Black Height}) \leq 2 \log_2(n + 1)$.

Optional Theory Problems: Bottleneck Paths

Problem 1: Modified Dijkstra for Bottleneck Paths

Goal: Compute a path where the maximum edge weight (bottleneck) is minimized. Time: $O(m \log n)$.

- **Algorithm:** Use Dijkstra's algorithm but modify the relaxation step.

- **Relaxation:** Instead of $d[v] = \min(d[v], d[u] + w(u, v))$, use:

$$d[v] = \min(d[v], \max(d[u], w(u, v)))$$

- **Explanation:** $d[u]$ represents the minimum bottleneck capacity to reach u . To extend this path to v , the new bottleneck is the larger of the previous bottleneck ($d[u]$) or the new edge itself ($w(u, v)$).

- **Complexity:** We use the same priority queue structure as standard Dijkstra. The number of operations is identical. Thus, $O(m \log n)$.

Problem 2: Undirected Graphs in Linear Time

Goal: Compute minimum-bottleneck path in undirected graph in $O(m)$.

- **Insight:** We can use a deterministic median-finding algorithm.

- **Algorithm (Sketch):**

1. Find the median edge weight w_{med} of the current edges ($O(m)$).
 2. Consider the subgraph G_{low} containing only edges with weight $\leq w_{med}$.
 3. Run BFS/DFS on G_{low} to see if s and t are in the same connected component ($O(m)$).
 4. **If Connected:** The optimal bottleneck is $\leq w_{med}$. Discard all edges $> w_{med}$ and recurse on G_{low} .
 5. **If Not Connected:** The optimal bottleneck is $> w_{med}$. Contract the connected components of G_{low} into super-nodes and recurse using the edges $> w_{med}$.
- **Complexity:** In each step, we reduce the number of edges by half or contract vertices significantly. The recurrence is $T(m) = T(m/2) + O(m)$, which solves to $O(m)$.

Problem 3: Directed Graphs

Goal: Can we do faster than $O(m \log n)$ for directed graphs?

- **Answer:** No (or at least, not using the simple contraction method).
- **Reasoning:** The linear-time strategy for undirected graphs relies on **contracting** connected components to reduce the problem size when the bottleneck is in the "upper half" of edge weights.
- In directed graphs, reachability is not symmetric. If s cannot reach t using only "light" edges, we cannot simply contract the components, because a "light" edge might still be required to bridge two "heavy" edges in a valid path.
- Therefore, we typically cannot beat the $O(m \log n)$ bound of the modified Dijkstra (or a bottleneck-sort based approach) easily.